

Custom Network Proxy Server (HTTP + HTTPS CONNECT)

Project Report

1. Abstract

This project implements a custom forward proxy server using Python socket programming. The proxy supports HTTP request forwarding, concurrent multi-client handling via a thread-per-connection model, domain-based blocking using a configurable blocklist, structured logging, and HTTPS support using CONNECT tunneling.

2. Objectives

- Implement a forward proxy using TCP sockets (HTTP/1.1).
- Handle multiple clients concurrently using threads.
- Parse HTTP requests to extract method, host, port, and path.
- Forward requests to destination servers and relay responses back.
- Implement rule-based filtering (blocklist) with 403 responses.
- Maintain logs for allowed/blocked requests and transferred byte counts.
- Support HTTPS through CONNECT tunneling (optional extension).

3. System Architecture

The proxy runs as a TCP server listening on 127.0.0.1:8888. For each incoming client connection, a dedicated thread is spawned to handle the session.

3.1 HTTP Forwarding Flow

1. Accept client TCP connection (browser/curl).
2. Receive request headers until CRLF CRLF (\r\n\r\n).
3. Parse request line and headers (method, target, HTTP version, Host header).
4. Resolve destination host/port/path from absolute-URI or Host + relative-path.
5. Apply blocklist policy. If blocked, respond with 403 Forbidden.
6. Open TCP connection to destination host:port.
7. Rewrite request line to server-form: METHOD /path HTTP/1.1.
8. Forward headers (+ body if Content-Length exists).
9. Relay the destination response back to the client until the server closes.
10. Log the result (ALLOW/BLOCKED/ERROR) and bytes transferred.

3.2 HTTPS CONNECT Tunneling Flow

11. Client sends CONNECT host:port HTTP/1.1 to the proxy.
12. Proxy parses the CONNECT target host and port (typically 443).
13. Apply blocklist policy to the CONNECT host. If blocked, respond 403 Forbidden.
14. Proxy opens a TCP connection to host:port.
15. Proxy replies: HTTP/1.1 200 Connection Established.
16. Create a transparent bidirectional tunnel (proxy does not decrypt TLS).
17. Log tunnel bytes in both directions (client→server and server→client).

4. Implementation Details

4.1 Concurrency Model

A thread-per-connection model is used. The main thread accepts connections and spawns a daemon thread for each client session. This approach is simple and satisfies the requirement to handle multiple clients.

4.2 Request Parsing

The proxy reads until `\r\n\r\n` to obtain complete HTTP headers. Headers are parsed into a dictionary using lowercase keys. The proxy supports both absolute-URI form (common when using a proxy) and relative-path form using the Host header.

4.3 Forwarding and Relay

For HTTP requests, the proxy connects to the destination server and forwards a rewritten request whose request line contains only the path. Responses are streamed back to the client in chunks, avoiding full buffering in memory. Total relayed bytes are recorded for logging.

4.4 Blocklist Filtering

Blocked domains are defined in config/blocked.txt. Matching is case-insensitive and includes parent-domain blocking (blocking example.com also blocks www.example.com). Blocked requests return HTTP/1.1 403 Forbidden and are logged.

4.5 Logging

Each request generates a log entry containing UTC timestamp, decision (ALLOW/BLOCKED/ERROR), client address, target, and byte counts. CONNECT tunnels log c2s and s2c byte counts.

5. Testing and Results

The proxy was tested using curl with explicit proxy settings. The following tests were executed:

Test	Command	Expected Result
------	---------	-----------------

HTTP Allow	<code>curl -v http://example.com -x http://127.0.0.1:8888</code>	200 OK, HTML returned, ALLOW log with bytes
HTTP Block	<code>curl -v http://example.com -x http://127.0.0.1:8888</code> (example.com in blocked.txt)	403 Forbidden, BLOCKED log
HTTPS Allow	<code>curl -v https://example.com -x http://127.0.0.1:8888</code>	200 Connection Established, then 200 OK, CONNECT log with c2s/s2c
HTTPS Block	<code>curl -v https://example.com -x http://127.0.0.1:8888</code> (example.com in blocked.txt)	403 Forbidden, CONNECT tunnel fails, BLOCKED CONNECT log

Screenshots of outputs and corresponding logs are attached in the Appendix.

6. Limitations

- Chunked request bodies are not fully supported (request-body handling assumes Content-Length).
- No caching layer is implemented.
- No authentication or user access control is implemented.
- Designed for HTTP/1.1 over TCP; does not support HTTP/2 or QUIC.

7. Future Improvements

- Add LRU-based HTTP response caching with validation headers (ETag/Last-Modified).
- Use a thread pool or async I/O for better scalability.
- Add optional authentication and access control lists (ACLs).
- Extend HTTP edge-case handling (chunked requests, pipelining).

Appendix A: Evidence (Screenshots)

Here are the screenshots of tests.

- HTTP Allow: curl output (200 OK + HTML).

```
D:\repo>curl -v http://example.com -x http://127.0.0.1:8888
* Trying 127.0.0.1:8888...
* Established connection to 127.0.0.1 (127.0.0.1 port 8888) from 127.0.0.1 port 58052
* using HTTP/1.x
> GET http://example.com/ HTTP/1.1
> Host: example.com
> User-Agent: curl/8.16.0
> Accept: */*
> Proxy-Connection: Keep-Alive
>
* Request completely sent off
< HTTP/1.1 200 OK
< Date: Thu, 08 Jan 2026 12:00:04 GMT
< Content-Type: text/html
< Transfer-Encoding: chunked
< Connection: close
< CF-RAY: 9bab8849981345af-DEL
< Last-Modified: Sat, 03 Jan 2026 05:43:21 GMT
< Allow: GET, HEAD
< Accept-Ranges: bytes
< Age: 5834
< cf-cache-status: HIT
< Server: cloudflare
<
<!doctype html><html lang="en"><head><title>Example Domain</title><meta name="viewport" content="width=device-width, initial-scale=1"><style>body{background:#eee; width:60vw; margin:15vh auto; font-family:system-ui,sans-serif}h1{font-size:1.5em}div{ opacity:0.8}>a:link,>a:visited{color:#348}</style></head><body><div><h1>Example Domain</h1><p>This domain is for use in documentation examples without needing permission. Avoid use in operations.<p><a href="https://iana.org/domains/example">Learn more</a></div></body></html>
* shutting down connection #0
```

- HTTP Block: curl output (403 Forbidden).

```
D:\repo>curl -v http://example.com -x http://127.0.0.1:8888
* Trying 127.0.0.1:8888...
* Established connection to 127.0.0.1 (127.0.0.1 port 8888) from 127.0.0.1 port 58056
* using HTTP/1.x
> GET http://example.com/ HTTP/1.1
> Host: example.com
> User-Agent: curl/8.16.0
> Accept: */*
> Proxy-Connection: Keep-Alive
>
< HTTP/1.1 403 Forbidden
< Connection: close
< Content-Type: text/plain
< Content-Length: 33
<
403 Forbidden (Blocked by proxy)
* we are done reading and this is set to close, stop send
* abort upload
* shutting down connection #0
```

- HTTPS Allow: curl output showing 200 Connection Established and final 200 OK.

```
D:\repo>curl -v https://example.com -x http://127.0.0.1:8888
* Trying 127.0.0.1:8888...
* CONNECT: no ALPN negotiated
* allocate connect buffer
* Establish HTTP proxy tunnel to example.com:443
> CONNECT example.com:443 HTTP/1.1
> Host: example.com:443
> User-Agent: curl/8.16.0
> Proxy-Connection: Keep-Alive
>
< HTTP/1.1 200 Connection Established
<
* CONNECT phase completed
* CONNECT tunnel established, response 200
* schannel: disabled automatic use of client certificate
* ALPN: curl offers http/1.1
* ALPN: server accepted http/1.1
* Established connection to 127.0.0.1 (127.0.0.1 port 8888) from 127.0.0.1 port 58058
* using HTTP/1.x
> GET / HTTP/1.1
> Host: example.com
> User-Agent: curl/8.16.0
> Accept: */*
>
* Request completely sent off
* schannel: remote party requests renegotiation
* schannel: renegotiating SSL/TLS connection
* schannel: SSL/TLS connection renegotiated
< HTTP/1.1 200 OK
Date: Thu, 08 Jan 2026 12:00:38 GMT
< Content-Type: text/html
< Transfer-Encoding: chunked
< Connection: keep-alive
< CF-RAY: 9bab89216db0c66-DEL
< last-modified: Sat, 03 Jan 2026 05:43:21 GMT
< allow: GET, HEAD
< Accept-Ranges: bytes
< Age: 6722
< cf-cache-status: HIT
< Server: cloudflare
<
<!doctype html><html lang="en"><head><title>Example Domain</title><meta name="viewport" content="width=device-width, initial-scale=1"><style>body{background:#eee; width:60vw; margin:15vh auto; font-family:systein-ui,sans-serif}</style></head><body><div><h1>Example Domain</h1><p>This domain is for use in documentation examples without needing permission. Avoid use in operations.<p><a href="https://iana.org/domains/example">Learn more</a></p></div></body></html>
* Connection #0 to host 127.0.0.1:8888 left intact
```

- HTTPS Block: curl output showing CONNECT blocked (403).

```
D:\repo>curl -v https://example.com -x http://127.0.0.1:8888
* Trying 127.0.0.1:8888...
* CONNECT: no ALPN negotiated
* allocate connect buffer
* Establish HTTP proxy tunnel to example.com:443
> CONNECT example.com:443 HTTP/1.1
> Host: example.com:443
> User-Agent: curl/8.16.0
> Proxy-Connection: Keep-Alive
>
< HTTP/1.1 403 Forbidden
< Connection: close
< Content-Type: text/plain
< Content-Length: 33
<
* CONNECT tunnel failed, response 403
* closing connection #0
curl: (56) CONNECT tunnel failed, response 403
```

- logs/proxy.log showing ALLOW/BLOCKED entries for HTTP and CONNECT (including bytes/c2s/s2c).

2026-01-08 12:00:04 UTC ALLOW 127.0.0.1:58052 -> GET example.com:80/ bytes=822
2026-01-08 12:00:22 UTC BLOCKED 127.0.0.1:58056 -> GET example.com:80/
2026-01-08 12:00:39 UTC ALLOW 127.0.0.1:58058 -> CONNECT example.com:443 c2s=659 s2c=5699
2026-01-08 12:00:56 UTC BLOCKED 127.0.0.1:58060 -> CONNECT example.com:443